


MAXIMUM ENTROPY MODEL

Yongmei Tan
ymtan@bupt.edu.cn

History

- The concept of Maximum Entropy can be traced back along multiple threads to Biblical times.
- Introduced to NLP area by Berger et. Al. (1996).
- Used in many NLP tasks: MT, Tagging, Parsing, PP attachment, LM, ...
- 1981年，托宾的“投资组合理论”使他获得了当年的诺贝尔经济学奖。



- 老头晒太阳

- I see a girl with a telescope



Modeling

The basic idea

- Goal: estimate p
- Choose p with maximum entropy (or “uncertainty”) subject to the constraints (or “evidence”).

$$H(p) = - \sum_{x \in A \times B} p(x) \log p(x)$$

$$x = (a, b), \quad \text{where} \quad a \in A \wedge b \in B$$

Setting

- From training data, collect (a, b) pairs:
 - a: thing to be predicted (e.g., a class in a classification problem)
 - b: the context
 - E.g., POS tagging:
 - a=NN
 - b=the words in a window and previous two tags
- Learn the prob of each (a, b): $p(a, b)$

Features in POS tagging (Ratnaparkhi, 1996)

Condition	Features	
w_i is not rare	$w_i = X$	$\& t_i = T$
w_i is rare	X is prefix of w_i , $ X \leq 4$	$\& t_i = T$
	X is suffix of w_i , $ X \leq 4$	$\& t_i = T$
	w_i contains number	$\& t_i = T$
	w_i contains uppercase character	$\& t_i = T$
	w_i contains hyphen	$\& t_i = T$
$\forall w_j$	$t_{i-1} = X$	$\& t_i = T$
	$t_{i-2}t_{i-1} = XY$	$\& t_i = T$
	$w_{j-1} = X$	$\& t_i = T$
	$w_{j-2} = X$	$\& t_i = T$
	$w_{j+1} = X$	$\& t_i = T$
	$w_{j+2} = X$	$\& t_i = T$

context (history)

allowable classes



Maximum Entropy

- Why maximum entropy?
- Maximize entropy = Minimize commitment
- Model all that is known and assume nothing about what is unknown.
 - Model all that is known: satisfy a set of constraints that must hold
 - Assume nothing about what is unknown: choose the most "uniform" distribution
→ choose the one with maximum entropy



COIN-FLIP EXAMPLE

- Toss a coin: $p(H)=p1$, $p(T)=p2$.
- Constraint: $p1 + p2 = 1$
- Question: what's your estimation of $p=(p1, p2)$?
- Answer: choose the p that maximizes $H(p)$


$$H(p) = -\sum p(x) \log p(x)$$



COIN-FLIP EXAMPLE (cont)



$$p_1 + p_2 = 1$$



$$p_1 + p_2 = 1.0, \quad p_1 = 0.3$$

An MT example(Berger et. al., 1996)

Possible translation for the word "in" is:

{dans, en, à, au cours de, pendant}

Constraint:

$$p(\textit{dans}) + p(\textit{en}) + p(\textit{à}) + p(\textit{au cours de}) + p(\textit{pendant}) = 1$$

Intuitive answer

$$p(\textit{dans}) = 1/5$$

$$p(\textit{en}) = 1/5$$

$$p(\textit{à}) = 1/5$$

$$p(\textit{au cours de}) = 1/5$$

$$p(\textit{pendant}) = 1/5$$

Constraints:

$$p(\textit{dans}) + p(\textit{en}) = 3/10$$

$$p(\textit{dans}) + p(\textit{en}) + p(\textit{\grave{a}}) + p(\textit{au cours de}) + p(\textit{pendant}) = 1$$

Intuitive answer:

$$p(\textit{dans}) = 3/20$$

$$p(\textit{en}) = 3/20$$

$$p(\textit{\grave{a}}) = 7/30$$

$$p(\textit{au cours de}) = 7/30$$

$$p(\textit{pendant}) = 7/30$$

Constraints:

$$p(\textit{dans}) + p(\textit{en}) = 3/10$$

$$p(\textit{dans}) + p(\textit{en}) + p(\textit{\`a}) + p(\textit{au cours de}) + p(\textit{pendant}) = 1$$

$$p(\textit{dans}) + p(\textit{\`a}) = 1/2$$

Intuitive answer:

??

POS tagging example (Klein and Manning, 2003)

Lets say we have the following event space:

NN	NNS	NNP	NNPS	VBZ	VBD
----	-----	-----	------	-----	-----

... and the following empirical data:

3	5	11	13	3	1
---	---	----	----	---	---

Maximize H:

$1/e$	$1/e$	$1/e$	$1/e$	$1/e$	$1/e$
-------	-------	-------	-------	-------	-------

... want probabilities: $\sum E[\text{NN}, \text{NNS}, \text{NNP}, \text{NNPS}, \text{VBZ}, \text{VBD}] = 1$

$1/6$	$1/6$	$1/6$	$1/6$	$1/6$	$1/6$
-------	-------	-------	-------	-------	-------

N^* are more common than V^* , so we add the feature $f_N = \{\text{NN}, \text{NNS}, \text{NNP}, \text{NNPS}\}$, with $\mathbb{E}[f_N] = 32/36$

NN	NNS	NNP	NNPS	VBZ	VBD
8/36	8/36	8/36	8/36	2/36	2/36

... and proper nouns are more frequent than common nouns, so we add $f_P = \{\text{NNP}, \text{NNPS}\}$, with $\mathbb{E}[f_P] = 24/36$

4/36	4/36	12/36	12/36	2/36	2/36
------	------	-------	-------	------	------

... we could keep refining the models, e.g. by adding a feature to distinguish singular vs. plural nouns, or verb types.

- 例6.1 假设随机变量 X 有5个取值 $\{A, B, C, D, E\}$, 要估计取各个值的概率 $P(A), P(B), P(C), P(D), P(E)$.

- $P(A) + P(B) + P(C) + P(D) + P(E) = 1$

- 从一些先验知识中得到对概率值的约束条件
 $P(A) + P(B) = 3/10$

- $P(A) + P(B) + P(C) + P(D) + P(E) = 1$

- 还有第三个约束条件

- $P(A) + P(C) = 1/2$

- $P(A) + P(B) = 3/10$

- $P(A) + P(B) + P(C) + P(D) + P(E) = 1$

Modeling the problem

- Objective function: $H(p)$
- Goal: Among all the distributions that satisfy the constraints, choose the one, p^* , that maximizes $H(p)$.

$$p^* = \arg \max_{p \in P} H(p)$$

- Question: How to represent constraints?

FEATURES

- Feature (feature function, Indicator function) is a binary-valued function on events:

$$f_j : \varepsilon \rightarrow \{0,1\}, \quad \varepsilon = A \times B$$

- A : the set of possible classes (e.g., tags in POS tagging)
- B : space of contexts (e.g., neighboring words/ tags in POS tagging)
- E.g.:

$$f_j(a,b) = \begin{cases} 1 & \text{if } a = DET \text{ \& } curWord(b) = "that" \\ 0 & \text{o.w.} \end{cases}$$

SOME NOTATIONS

Finite training sample of events:

$$S$$

Observed probability of x in S :

$$\tilde{p}(x)$$

The model P 's probability of x :

$$p(x)$$

The j^{th} feature:

$$f_j$$

Observed expectation of f_j :
(empirical count of f_j)

$$E_{\tilde{p}} f_j = \sum_{x \in \mathcal{E}} \tilde{p}(x) f_j(x)$$

Model expectation of f_j :

$$E_p f_j = \sum_{x \in \mathcal{E}} p(x) f_j(x)$$

CONSTRAINTS

- Model's feature expectation = observed feature expectation

$$E_p f_j = E_{\tilde{p}} f_j$$


- How to calculate $E_{\tilde{p}} f_j$?

$$E_{\tilde{p}} f_j = \sum_{x \in \mathcal{E}} \tilde{p}(x) f_j(x) = \frac{\sum_{l=1}^N f_j(x_l)}{N}$$

$$f_j(a, b) = \begin{cases} 1 & \text{if } a = DET \text{ \& } curWord(b) = "that" \\ 0 & \text{o.w.} \end{cases}$$

TRAINING DATA → OBSERVED EVENTS

<i>Word:</i>	the	stories	about	well-heeled	communities	and	developers
<i>Tag:</i>	DT	NNS	IN	JJ	NNS	CC	NNS
<i>Position:</i>	1	2	3	4	5	6	7



$w_i = \text{about}$ $\& t_i = \text{IN}$
 $w_{i-1} = \text{stories}$ $\& t_i = \text{IN}$
 $w_{i-2} = \text{the}$ $\& t_i = \text{IN}$
 $w_{i+1} = \text{well-heeled}$ $\& t_i = \text{IN}$
 $w_{i+2} = \text{communities}$ $\& t_i = \text{IN}$
 $t_{i-1} = \text{NNS}$ $\& t_i = \text{IN}$
 $t_{i-2}t_{i-1} = \text{DT NNS}$ $\& t_i = \text{IN}$

RESTATING THE PROBLEM

The task: find p^* s.t.

$$p^* = \arg \max_{p \in P} H(p)$$

where

$$P = \{p \mid E_p f_j = E_{\tilde{p}} f_j, j = \{1, \dots, k\}\}$$

Objective function: $H(p)$

Constraints:

$$\{E_p f_j = E_{\tilde{p}} f_j = d_j, j = \{1, \dots, k\}\}$$

$$\sum_x p(x) = 1$$



Add a feature $f_0(a, b) = 1 \quad \forall a, b$

$$E_p f_0 = E_{\tilde{p}} f_0 = 1$$

Questions

- What is the definition of p^* ?
- What is the form of p^* ?
- How to find p^* ?

What is the definition of p^* ?

- 假设满足所有约束条件的模型集合为 $\mathcal{C}$:

$$\mathcal{C} \equiv \{P \in \mathcal{P} | E_P(f_i) = E_{\tilde{P}}(f_i), i = 1, 2, \dots, n\}$$

$\mathcal{P}$ 为没有任何约束条件时所有概率分布的集合

$P(Y|X)$ 为条件概率分布

集合 $\mathcal{P}$ 中的一个元素, 即 $P(Y|X) \in \mathcal{P}$

可得出 $\mathcal{C}$ 为 $\mathcal{P}$ 的子集, 即 f_i 为第 i 个特征函数, 一共有 n 个特征函数。

What is the definition of p^* ?

定义在条件概率分布 $P(Y|X)$ 上的条件熵为：

$$H(P) = - \sum_{x,y} \tilde{P}(x) P(y|x) \log P(y|x)$$

$\tilde{P}(x)$ 为边缘分布 $P(x)$ 的经验分布，对数为自然对数。条件熵的最小值为0，最大值为 $\log|Y|$ ，其中 Y 为其可能取值的个数。

则最大熵模型即为在集合 $\mathcal{C}$ 中找出满足以下条件的概率分布：

$$P^* = \operatorname{argmax}_{P \in \mathcal{C}} H(P)$$

$P^* \in \mathcal{C}$ ， $H(P)$ 为一个概率分布的条件熵。

$$p^* = \arg \max_{p \in P} \sum_{(x,y)} p(y | x) \tilde{p}(x) \log \frac{1}{p(y | x)}$$

$$P = \left\{ p(y | x) \left| \begin{array}{l} \forall f_i : \\ \forall x : \end{array} \right. \begin{array}{l} \sum_{(x,y)} p(y | x) \tilde{p}(x) f_i(x, y) \\ \sum_{(x,y)} \tilde{p}(x, y) f_i(x, y) \\ \sum_y p(y | x) = 1 \end{array} \right\} =$$

定义条件熵 $H(y|x) = - \sum_{(x,y) \in Z} p(y,x) \log p(y|x)$

模型目标 $p^*(y|x) = \arg \max_{p(y|x) \in P} H(y|x)$

定义特征函数 $f_i(x, y) \in \{0, 1\} \quad i = 1, 2, \dots, m$

约束条件 $\sum_{y \in Y} p(y|x) = 1$

$E(f_i) = \tilde{E}(f_i) \quad i = 1, 2, \dots, m \quad N = |T|$ ⁽²⁾

$\tilde{E}(f_i) = \sum_{(x,y) \in Z} \tilde{p}(x, y) f_i(x, y) = \frac{1}{N} \sum_{(x,y) \in T} f_i(x, y)$

$E(f_i) = \sum_{(x,y) \in Z} p(x, y) f_i(x, y) = \sum_{(x,y) \in Z} p(x) p(y|x) f_i(x, y)$

WHAT IS THE FORM OF P^* ?

(RATNAPARKHI, 1997)

$$P = \{p \mid E_p f_j = E_{\tilde{p}} f_j, j = \{1, \dots, k\}\}$$

$$Q = \{p \mid p(x) = \pi \prod_{j=1}^k \alpha_j^{f_j(x)}, \alpha_j > 0\}$$

Theorem: if $p^* \in P \cap Q$ then $p^* = \arg \max_{p \in P} H(p)$
Furthermore, p^* is unique.

How to find p^* ?

- 最大熵模型的学习转化为约束最优化问题：

$$\begin{aligned} \max_{P \in \mathcal{C}} \quad & H(P) = - \sum_{x,y} \tilde{P}(x) P(y|x) \log P(y|x) \\ \text{s.t.} \quad & \forall x, y: P(y|x) \geq 0 \\ & \sum_y P(y|x) = 1 \text{ for all } x \\ & E_P(f_i) = E_{\tilde{P}}(f_i), \quad i = 1, 2, \dots, n \end{aligned}$$

How to find p^* ?

- 按照求解最优化问题的习惯，将求最大值问题转化为求解最小值的问题：

$$\min_{p \in \mathcal{C}} -H(P) = \sum_{x,y} \tilde{P}(x) P(y|x) \log P(y|x)$$

$$\text{s.t. } \forall x, y: P(y|x) \geq 0$$

$$\sum_y P(y|x) = 1 \text{ for all } x$$

$$E_P(f_i) = E_{\tilde{P}}(f_i), \quad i = 1, 2, \dots, n$$

拉格朗日函数 $L(p, \lambda)$

$$\begin{aligned} L(P, \lambda) &= -H(P) + \lambda_0 \left(1 - \sum_y P(y|x) \right) \\ &\quad + \sum_{i=1}^n \lambda_i (E_{\tilde{P}}(f_i) - E_P(f_i)) \\ &= \sum_{x,y} \tilde{P}(x) P(y|x) \log P(y/x) + \lambda_0 \left(1 - \sum_y P(y|x) \right) \\ &\quad + \sum_{i=1}^n \lambda_i \left(\sum_{x,y} \tilde{P}(x, y) f_i(x, y) - \sum_{x,y} \tilde{P}(x) P(y|x) f_i(x, y) \right) \end{aligned}$$

$$\begin{aligned}
& \frac{\partial L(P, \boldsymbol{\Lambda})}{\partial P(y|x)} \\
&= \sum_{x,y} \tilde{P}(x)(\log P(y|x) + 1) - \sum_y \boldsymbol{\Lambda}_0 \\
&\quad - \sum_{x,y} \left(\tilde{P}(x) \sum_{i=1}^n \boldsymbol{\Lambda}_i f_i(x, y) \right) \\
&= \sum_{x,y} \tilde{P}(x)(\log P(y|x) + 1 - \boldsymbol{\Lambda}_0 - \\
&\quad \sum_{i=1}^n \boldsymbol{\Lambda}_i f_i(x, y))
\end{aligned}$$

令上述偏导数等于0，在 $\tilde{P}(x) > 0$ 的情况下

$$P(y|x) = \exp(\sum_{i=1}^n \lambda_i f_i(x, y) + \lambda_0 - 1) = \frac{\exp(\sum_{i=1}^n \lambda_i f_i(x, y))}{\exp(1 - \lambda_0)}$$

任意 x 有 $\sum_y p(y|x) = 1$ ，令 $1/\exp(1 - \lambda_0)$ 为常数 c

$$\sum_y c e^{\sum_i \lambda_i f_i(x, y)} = 1$$

解得：

$$c = \frac{1}{\sum_y e^{\sum_i \lambda_i f_i(x, y)}}$$

令

$$Z(x) = \sum_y e^{\sum_i \lambda_i f_i(x, y)}$$

$$P_{\lambda}(y|x) = \frac{1}{Z_{\lambda}(x)} \exp(\sum_{i=1}^n \lambda_i f_i(x, y))$$

$$Z_{\lambda}(x) = \sum_y \exp\left(\sum_{i=1}^n \lambda_i f_i(x, y)\right)$$

- 例6.2 假设随机变量 X 有5个取值 $\{A, B, C, D, E\}$, 要估计取各个值的概率 $P(A), P(B), P(C), P(D), P(E)$. 从一些先验知识中得到对概率值的约束条件 $P(A)+P(B)=3/10$ 。

例 6.2 学习例 6.1 中的最大熵模型.

解 为了方便, 分别以 y_1, y_2, y_3, y_4, y_5 表示 A, B, C, D 和 E , 于是最大熵模型学习的最优化问题是

$$\begin{aligned} \min \quad & -H(P) = \sum_{i=1}^5 P(y_i) \log P(y_i) \\ \text{s.t.} \quad & P(y_1) + P(y_2) = \tilde{P}(y_1) + \tilde{P}(y_2) = \frac{3}{10} \\ & \sum_{i=1}^5 P(y_i) = \sum_{i=1}^5 \tilde{P}(y_i) = 1 \end{aligned}$$

引进拉格朗日乘子 w_0, w_1 , 定义拉格朗日函数

$$L(P, w) = \sum_{i=1}^5 P(y_i) \log P(y_i) + w_1 \left(P(y_1) + P(y_2) - \frac{3}{10} \right) + w_0 \left(\sum_{i=1}^5 P(y_i) - 1 \right)$$

根据拉格朗日对偶性，可以通过求解对偶最优化问题得到原始最优化问题的解，所以求解

$$\max_w \min_P L(P, w)$$

首先求解 $L(P, w)$ 关于 P 的极小化问题。为此，固定 w_0, w_1 ，求偏导数：

$$\frac{\partial L(P, w)}{\partial P(y_1)} = 1 + \log P(y_1) + w_1 + w_0$$

$$\frac{\partial L(P, w)}{\partial P(y_2)} = 1 + \log P(y_2) + w_1 + w_0$$

$$\frac{\partial L(P, w)}{\partial P(y_3)} = 1 + \log P(y_3) + w_0$$

$$\frac{\partial L(P, w)}{\partial P(y_4)} = 1 + \log P(y_4) + w_0$$

$$\frac{\partial L(P, w)}{\partial P(y_5)} = 1 + \log P(y_5) + w_0$$

令各偏导数等于 0，解得

$$P(y_1) = P(y_2) = e^{-w_1 - w_0 - 1}$$

$$P(y_3) = P(y_4) = P(y_5) = e^{-w_0 - 1}$$

于是，

$$\min_P L(P, w) = L(P_w, w) = -2e^{-w_1 - w_0 - 1} - 3e^{-w_0 - 1} - \frac{3}{10}w_1 - w_0$$

于是,

$$\min_P L(P, w) = L(P_w, w) = -2e^{-w_1 - w_0 - 1} - 3e^{-w_0 - 1} - \frac{3}{10}w_1 - w_0$$

再求解 $L(P_w, w)$ 关于 w 的极大化问题:

$$\max_w L(P_w, w) = -2e^{-w_1 - w_0 - 1} - 3e^{-w_0 - 1} - \frac{3}{10}w_1 - w_0$$

分别求 $L(P_w, w)$ 对 w_0, w_1 的偏导数并令其为 0, 得到

$$e^{-w_1 - w_0 - 1} = \frac{3}{20}$$

$$e^{-w_0 - 1} = \frac{7}{30}$$

于是得到所要求的概率分布为

$$P(y_1) = P(y_2) = \frac{3}{20}$$

$$P(y_3) = P(y_4) = P(y_5) = \frac{7}{30}$$

对偶函数的极大化等价于
最大熵的极大似然估计

似然函数与极大似然估计

- 设随机样本 $X_1, X_2, \dots, X_n$ 来自分布密度（或概率函数）为 $f(X, \theta)$ 的总体，其中 θ 是未知参数，则

$$L(\theta; X_1, \dots, X_n) = f(X_1, \theta) \dots f(X_n, \theta)$$

称为样本 $(X_1, X_2, \dots, X_n)$ 的似然函数（likelihood function），简记为 $L(\theta)$ 。

似然函数与极大似然估计

- 使似然函数 $L(\theta)$ 取最大值的 $\hat{\theta}$ ，即

$$L(\hat{\theta}) = \sup_{\theta} L(\theta)$$

称为 θ 的极大似然估计 (maximum likelihood estimator, **MLE**)，记之为 $\hat{\theta}_{\text{MLE}}$ ，等价于求 $\hat{\theta}$ 。

- 极大似然估计是在给定模型（含有未知参数）和数据集的情况下，估计出最佳的模型中未知参数的值，这些参数的取值可以使模型实现对给定数据集最大程度的拟合。

似然函数与极大似然估计

- 使得

$$\ln L(\hat{\theta}) = \max_{\theta} \ln L(\theta)$$

可得

$$\frac{d \ln L(\theta)}{d \theta} = 0$$

称为似然方程。

- 当有多个参数时，上述似然方程变为方程组
$$\frac{\partial \ln L(\theta_1, \dots, \theta_k)}{\partial \theta_i} = 0, i \leq k$$

其中 k 为参数的个数。

- 例：给定一个硬币，此硬币有正面和反面两个不同的面，使用此硬币进行**10**次的抛硬币实验，最终得到的**10**次实验的结果为：正面，正面，反面，正面，正面，正面，反面，反面，正面，正面。令每次抛硬币结果为正面的概率为 p ，则根据极大似然估计方法，算出 p 的取值。

解：

由题意可知，每次抛硬币结果为反面的概率为 $1-p$ 。

由于每次抛硬币实验是相互独立的，所以得到上述实验结果的概率为：

$$P = pp(1-p)ppp(1-p)(1-p)pp = p^7(1-p)^3$$

极大似然估计就是要找出与样本的分布最接近的概率分布模型，因此要求上式取得最大值，即：

$$\max P = \max_{0 \leq p \leq 1} p^7(1-p)^3$$

对参数 p 进行求导，并令其求导之后的算式为0，即可求出上式最优解为 $p=0.7$ 。

- 例：考虑投掷一枚硬币的实验。假如已知投出的硬币正面朝上的概率是 $P_H = 0.5$ ，便可知投掷若干次后出现各种结果的可能性。

投两次都是正面朝上的概率是0.25

$$P_{HH} = 0.25$$

即给定“投两次都是正面朝上”的观测，则硬币正面朝上的概率为0.5的似然是

$$L(p_H = 0.5; HH) = P(HH | p_H = 0.5) = 0.25$$

如果考虑 $p_H = 0.6$ ，那么似然函数的值会变大

$$L(p_H = 0.6; HH) = 0.36.$$

如果参数的取值变成**0.6**的话，结果观测到连续两次正面朝上的概率要比假设**0.5** 时更大。

参数取成**0.6** 要比取成**0.5** 更有说服力，更为“合理”。

- 似然函数的重要性不是它的具体取值，而是当参数变化时函数到底变小还是变大。
- 对同一个似然函数，如果存在一个参数值，使得它的函数值达到最大的话，那么这个值就是最为“合理”的参数值。

最大熵模型中条件概率分布 $P(y|x)$ 对数似然函数为：

$$L_{\tilde{P}}(P) = \log \prod_{x,y} P(y|x)^{\tilde{P}(x,y)} = \sum_{x,y} \tilde{P}(x,y) \log P(y|x)$$

其中 $\tilde{P}(x,y)$ 为已知的从训练数据中得到的经验概率分布。

$$P_{\Lambda}(y|x) = \frac{1}{Z_{\Lambda}(x)} \exp(\sum_{i=1}^n \lambda_i f_i(x,y))$$

$$Z_{\Lambda}(x) = \sum_y \exp\left(\sum_{i=1}^n \lambda_i f_i(x,y)\right)$$

$$\begin{aligned} L_{\tilde{P}}(P_{\Lambda}) &= \sum_{x,y} \tilde{P}(x,y) \log P(y|x) \\ &= \sum_{x,y} \tilde{P}(x,y) \left(\sum_{i=1}^n \lambda_i f_i(x,y) - \log Z_{\Lambda}(x) \right) \end{aligned}$$

$$= \sum_{x,y} \tilde{P}(x,y) \sum_{i=1}^n \lambda_i f_i(x,y) - \sum_{x,y} \tilde{P}(x,y) \log Z_{\lambda}(x)$$

$$= \sum_{x,y} \tilde{P}(x,y) \sum_{i=1}^n \lambda_i f_i(x,y) - \sum_x \tilde{P}(x) \left(\sum_y P_{\lambda}(y|x) \right) \log Z_{\lambda}(x)$$

$$= \sum_{x,y} \tilde{P}(x,y) \sum_{i=1}^n \lambda_i f_i(x,y) - \sum_x \tilde{P}(x) \log Z_{\lambda}(x)$$

$$Z_{\lambda}(x) = \sum_y e^{\sum_i \lambda_i f_i(x,y)}, \text{ 并且 } \tilde{P}(x,y) = \tilde{P}(x)P(y|x)$$

$$\max_{P \in \mathcal{C}} \quad H(P) = -\sum_{x,y} \tilde{P}(x)P(y|x) \log P(y|x)$$

$$\text{s.t.} \quad E_P(f_i) = E_{\tilde{P}}(f_i), \quad i = 1, 2, \dots, n$$

$$\sum_y P(y|x) = 1$$

$$\min_{P \in \mathcal{C}} \quad -H(P) = \sum_{x,y} \tilde{P}(x)P(y|x) \log P(y|x)$$

$$\text{s.t.} \quad E_P(f_i) - E_{\tilde{P}}(f_i) = 0, \quad i = 1, 2, \dots, n$$

$$\sum_y P(y|x) = 1$$

$$L(P, w) \equiv -H(P) + w_0 \left(1 - \sum_y P(y|x) \right) + \sum_{i=1}^n w_i (E_{\tilde{P}}(f_i) - E_P(f_i))$$

$$= \sum_{x,y} \tilde{P}(x)P(y|x) \log P(y|x) + w_0 \left(1 - \sum_y P(y|x) \right)$$

$$+ \sum_{i=1}^n w_i \left(\sum_{x,y} \tilde{P}(x,y) f_i(x,y) - \sum_{x,y} \tilde{P}(x)P(y|x) f_i(x,y) \right)$$

$$\min_{P \in \mathcal{C}} \max_w L(P, w)$$

$$\max_w \min_{P \in \mathcal{C}} L(P, w)$$

$$\Psi(w) = \min_{P \in \mathcal{C}} L(P, w) = L(P_w, w)$$

$\Psi(w)$ 称为对偶函数. 同时, 将其解记作

$$P_w = \arg \min_{P \in \mathcal{C}} L(P, w) = P_w(y|x)$$

对于对偶函数 $\Psi(w)$

$$\begin{aligned}\Psi(\lambda) &= \sum_{x,y} \tilde{P}(x) P_{\lambda}(y|x) \log P_{\lambda}(y|x) \\ &\quad + \sum_{i=1}^n \lambda_i \left(\sum_{x,y} \tilde{P}(x,y) f_i(x,y) - \sum_{x,y} \tilde{P}(x) P_{\lambda}(y|x) f_i(x,y) \right) \\ &= \sum_{x,y} \tilde{P}(x,y) * \sum_{i=1}^n \lambda_i f_i(x,y) + \sum_{x,y} \tilde{P}(x) P_{\lambda}(y|x) \left(\log P_{\lambda}(y|x) - \sum_{i=1}^n \lambda_i f_i(x,y) \right) \\ &= \sum_{x,y} \tilde{P}(x,y) \sum_{i=1}^n \lambda_i f_i(x,y) - \sum_{x,y} \tilde{P}(x) P_{\lambda}(y|x) \log Z_{\lambda}(x) \\ &= \sum_{x,y} \tilde{P}(x,y) \sum_{i=1}^n \lambda_i f_i(x,y) - \sum_x \tilde{P}(x) \log Z_{\lambda}(x)\end{aligned}$$

其中最后一步用到了 $\sum_y P(y|x) = 1$ 。

可以得到：

$$\Psi(\lambda) = L_{\tilde{P}}(P_{\lambda})$$

Summary

- 最大熵模型学习中的对偶问题极大化等价于极大似然估计。
- 最大熵的解是符合数据集中数据分布的最优解，证明了最大熵模型的合理性。

THANK YOU!